Anusha Alangar

480-906-5917 | anushalngr@gmail.com | linkedin.com/in/anushaa51 | anushaa51.github.io

Education

Master of Science in Computer Science — *GPA: 4.0* Aug. 2023 – May 2025
Arizona State University *Tempe, AZ*

Skills

Languages & Tools: C/C++, Java, Kotlin, Python, Go, JavaScript, Gradle, Bazel, Jenkins, gRPC, Protobuf
Frameworks: Node.js, React, GraphQL, Next.js, LangChain, MCP, Spring Boot, Micronaut, FastAPI
Cloud: AWS EC2, EKS, S3, Lambda, Grafana, Elastic Search, Kibana, Docker, Kubernetes
Data: Kafka, Cassandra, Hadoop, PostgreSQL, Redis, MongoDB, ZooKeeper, InfluxDB, Prometheus

Experience

Research Assistant July 2025 – Present
VISA Research Lab - Arizona State University *Tempe, AZ*

- Researching social health recommendation agents using OpenAI models, focusing on adaptive personalization
- Designing solutions for long term memory context retrieval with open-source MCPs and Agentic RAG, leveraging LangChain for knowledge grounding and dialogue coherence

Software Engineering Intern June 2024 – Aug 2024
National Basketball Association (NBA) *New York City, NY*

- Worked on B2C features including Sports News, Ads, Catalogs, Scores and Subscriptions on NBA.com
- Performed Tailwind to CSS codebase migrations, abstracting serialization into separate packages
- Refactored core Python authentication APIs for League Pass subscriptions, and created a framework for granular control over ads serving logic in GraphQL

Software Engineer Aug 2020 – June 2023
Target Corporation *Bangalore, KA*

- Search team at Target is responsible for search and recommendation on Target.com, managing the entire lifecycle from search input and autocomplete, to query completion, normalization, and redirection
- Built and maintained enterprise grade cross-functional microservices touching Java, Kotlin, React, Docker, Kubernetes and Elasticsearch, along with associated observability and ops through Prometheus and Grafana
- Improved user query suggestion ranking by leveraging search prefix prioritization, fuzzy matching and look-up tries, leading to large improvements in peak search transaction throughput
- Productionized a Kafka based Tcin Sequence Generator, serving as an efficient product catalog engine that consolidates information from various sources
- Built pipelines to consolidate Kafka clusters and Redis cache in real-time for high availability during peak hours, and migrated sections of Target's tech stack from Java, Spring and MySQL to Kotlin, Micronaut and PostgreSQL
- Identified critical fields and re-indexed Target's fork of Metabase that interfaces with MongoDB, resulting in reduced time taken to visualize data from over a minute down to sub 1000ms

Projects

MidLLaMAI: Benchmarking suite for evaluating performance of compressed LLaMA models, focused on finding an optimal balance between different pruning and quantization techniques with varying model sizes

Crypticket: An offline capable cryptographic ticket generation and authentication platform using service workers and local storage caching. Built as a responsive PWA using React, utilizing EdDSA elliptic curve cryptography

Accentrix: A system built with TensorFlow for accent conversion and classification through mapping MFCCs vectors describing the short-term power spectrum of sound, including a React based web interface

LookUpBloodDB: Unified platform for locating the nearest blood bank with availability by interfacing with and consolidating blood donation databases across organizations, built with FastAPI and MySQL, exposed as a REST API

Awards & Honors

R&D Recognition Award | *Target Corporation* Oct. 2021

- Developed a high impact product catalog consolidator across Target's inventory in the enterprise search pipeline